

# Kubernetes

Manage a cluster of Linux containers in multi-cloud world, public, private or hybrid.

Luciano Faustino , Tiago Katcipis

# Agenda

- Introduction
- Why Kubernetes ?
- Concepts
- Deploying
- Service Management
- Monitoring

# Introduction

# What is Kubernetes?



**kubernetes**  
by Google™

- Kubernetes is an open-source platform for automating deployment, scaling, and operations of application containers across clusters of hosts
- Scale your applications on the fly
- Seamlessly roll out new features
- Optimize use of your hardware by using only the resources you need
- Kubernetes is: lean, portable, extensible, self-healing

# History

- The Kubernetes project was started by Google in 2014
- Kubernetes builds upon a decade and a half of experience that Google has with running production workloads at scale, combined with best-of-breed ideas and practices from the community
- Google's Borg system

[research.google.com/pubs/pub43438.html](https://research.google.com/pubs/pub43438.html) (<https://research.google.com/pubs/pub43438.html>)

[kubernetes.io/community/](http://kubernetes.io/community/) (<http://kubernetes.io/community/>)

# Why Kubernetes?

## Some problems solved


- Keep services running
- Stop services if necessary
- Schedule services around a cluster
- Load balance
- Service discovery
- Health checking
- Data management
- Support to autoscale
- Support to rolling updates
- Multi cloud

# Kubernetes Concepts



# Cluster

- A cluster is a set of physical or virtual machines and other infrastructure resources used by Kubernetes to run your applications



# Node

- Is a worker machine in Kubernetes, previously known as Minion
- Node may be a VM or physical machine, depending on the cluster
- Each node has the services necessary to run Pods and is managed by the master components

# Pods

- It is a set of containers
- They share all resources (context)
- Each pod has its own IP, that is shared through all containers inside it
- Models a logical host (that is why each one has its own IP)

# Labels

- key/value pairs that are attached to objects, such as pods.
- Used to organize and select groups of objects
- Labels enable users to map their own organizational structures onto system objects in a loosely coupled fashion, without requiring clients to store these mappings

Example labels:

```
"release" : "stable", "release" : "canary"
```

or

```
"environment" : "dev", "environment"* : "qa", "environment"* : "production"
```

# Label Selector

- Unlike names and UUIDs, labels do not provide uniqueness. In general, we expect many objects to carry the same label(s)
- Via a label selector, the client/user can identify a set of objects

Example labels selector:

```
environment = production
```

The former selects all resources with key equal to environment and value equal to production

```
environment in (production, qa)
```

The first example selects all resources with key equal to environment and value equal to production or qa

# Replication Controllers

- Ensures that a specified number of pod "replicas" are running at any one time
- If there are too many, it will kill some
- If there are too few, it will start more
- Replaces pods that are deleted or terminated for any reason, such as in the case of node failure
- Think of it similarly to a cluster aware process supervisor
- Useful to perform rolling updates

# Services

- Defines a logical set of Pods and a policy by which to access them
- Offers a virtual-IP-based bridge to Services which redirects to the backend Pods
- Service consumers always see the same IP, even if pods IPs changes
- Pods set that composes a service is built using a Label Selector

# Volumes

- The Kubernetes Volume abstraction (Persistent Volume)
- Kubernetes supports several types of Volumes: emptyDir, hostPath, gcePersistentDisk, awsElasticBlockStore, nfs, iscsi, glusterfs, rbd, gitRepo, secret, persistentVolumeClaim



# Secrets

- Objects of type secret are intended to hold sensitive information, such as passwords, OAuth tokens, and ssh keys
- Putting this information in a secret is safer and more flexible than putting it verbatim (text plain) in a pod definition or in a docker image

# Names

- Names are generally client-provided
- Only one object of a given kind can have a given name at a time
- Names are the used to refer to an object in a resource URL

```
GET /api/v1/pods/some-name
```

# Namespaces

- Intended for use in environments with many users spread across multiple teams, or projects
- A way to divide cluster resources between multiple uses
- Kubernetes starts with two initial namespaces: default, kube-system

# Annotations

- Like labels, annotations are key-value maps
- Possible information that could be recorded in annotations:

```
lightweight rollout tool metadata
```

or

```
build/release/image information (timestamps, release ids, git branch,  
PR numbers, image hashes, registry address, etc.)
```

or

```
phone/pager number(s) of person(s) responsible, or directory entry where  
that info could be found, such as a team website, etc.
```

# Deploying

- \* Provisioning cluster nodes
- \* Configuring a Kubernetes compatible network
- \* Deploying Kubernetes services and client tools
- \* Kubernetes can run on a range of platforms

[github.com/kubernetes/kubernetes/blob/master/docs/getting-started-guides/README.md#table-of-solutions](https://github.com/kubernetes/kubernetes/blob/master/docs/getting-started-guides/README.md#table-of-solutions) (<https://github.com/kubernetes/kubernetes/blob/master/docs/getting-started-guides/README.md#table-of-solutions>)

[guides/README.md#table-of-solutions](https://github.com/kubernetes/kubernetes/blob/master/docs/getting-started-guides/README.md#table-of-solutions))

# Operation System

## \* CoreOS

CoreOS is a minimal operating system that supports popular container systems out of the box. The operating system is designed to be operated in clusters.

[coreos.com](https://coreos.com) (<https://coreos.com>)

## Why?

- Cluster
- Minimal
- Fastpatch
- Bootstrap with cloud-config
- Container

# Automation Tools

## \* Cloud Machine

This is a Go Project that should be used to create a cloud environment. The app will create volumes and instance through AWS, although in the next future it'll be possible use other backends like Microsoft Azure, Google Cloud Platform, etc.

[github.com/NeowayLabs/cloud-machine](https://github.com/NeowayLabs/cloud-machine) (<https://github.com/NeowayLabs/cloud-machine>)

## Why?

- Simple
- Fast
- Without Bugs
- Write in Golang
- ????

# Kubernetes Control Plane (Master)

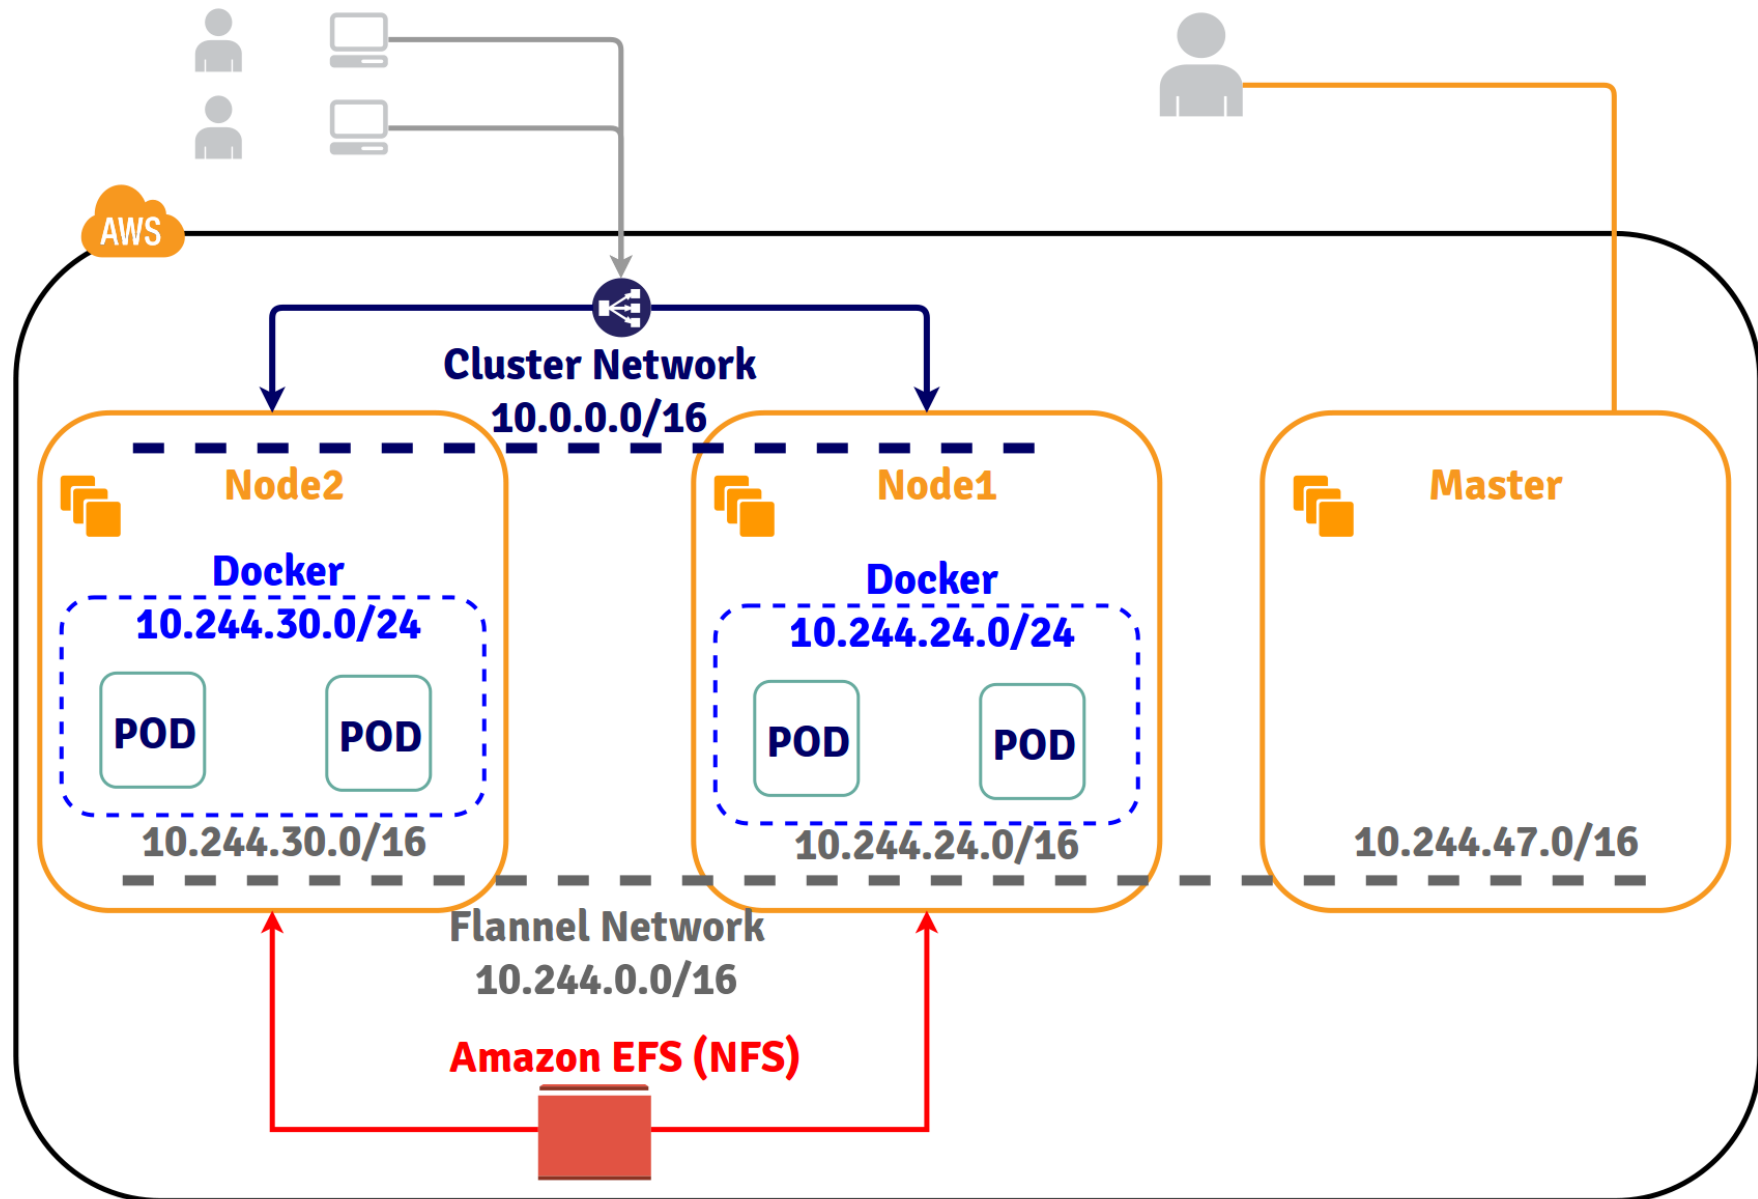
- **Kubernetes API Server** : The apiserver serves up the Kubernetes API
- **Scheduler** : Watches the apiserver for unscheduled pods and schedules them onto healthy nodes based on resource requirements
- **Controller Manager Server** : All other cluster-level functions
- **etcd** : A distributed consistent key-value store for shared configuration and service discovery. All persistent master state is stored in an instance of etcd
- **flannel** : Is a very simple overlay network. Provide SDN for Kubernetes
- **Docker Mirror** : Keep most of the image fetch traffic on your local network



# Kubernetes Node


- **Kubelet** : Manages pods and their containers, their images, their volumes, etc
- **Kube-Proxy** : Simple network proxy and load balancer
- **etcd** : A distributed consistent key-value store for shared configuration and service discovery. All persistent master state is stored in an instance of etcd
- **flannel** : Is a very simple overlay network. Provide SDN for Kubernetes
- **Docker** : Linux Container :-)

# Kubernetes Network



 Text

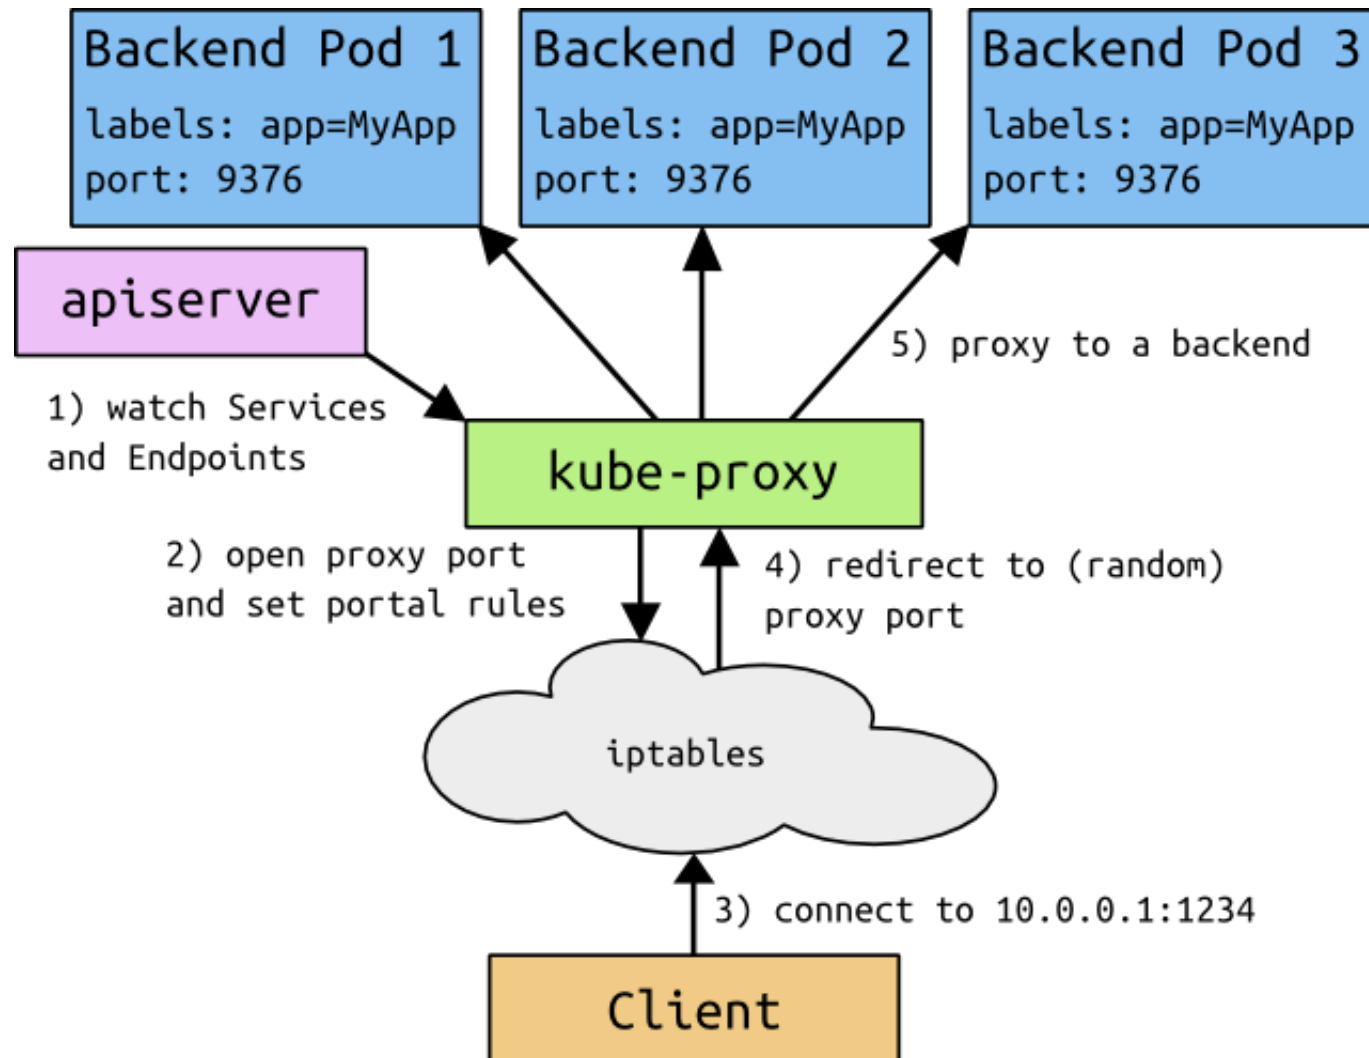
# Cluster Network



## Cluster Network



# Node Network Inside



# Service Management

# Running a service

```
kubectl -s http://{HOST}:{PORT} create -f k8s/rc/nginx.yaml
```

```
# cat k8s/rc/nginx.yaml
```

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: nginx
spec:
  replicas: 2
  selector:
    app: nginx
  template:
    metadata:
      name: nginx
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
          ports:
            - containerPort: 80
```



# Exposing a service

```
kubectl -s http://{HOST}:{PORT} create -f k8s/svc/nginx.yml
```

```
cat k8s/svc/nginx.yml
```

```
apiVersion: v1
kind: Service
metadata:
  labels:
    name: nginx-service
  name: nginx-service
spec:
  deprecatedPublicIPs:
    - 192.168.33.11
  ports:
    - port: 80 # the port that this service should serve on
      # the container on each pod to connect to, can be a name
      # (e.g. 'www') or a number (e.g. 80)
      targetPort: 80
      protocol: TCP
  # just like the selector in the replication controller,
  # but this time it identifies the set of pods to load balance
  # traffic to.
  selector:
    app: nginx
```

# Service Discovery

- Kubernetes supports 2 primary modes of finding a Service

## \* Environment variables

- When a Pod is run on a Node, the kubelet adds a set of environment variables for each active Service
- Supports both Docker links compatible variables and simpler **{SVCNAME}\_SERVICE\_HOST** and **{SVCNAME}\_SERVICE\_PORT** variables

For example, the Service "**redis-master**" which exposes TCP port 6379

```
REDIS_MASTER_SERVICE_HOST=10.0.0.11
REDIS_MASTER_SERVICE_PORT=6379
REDIS_MASTER_PORT=tcp://10.0.0.11:6379
REDIS_MASTER_PORT_6379_TCP=tcp://10.0.0.11:6379
REDIS_MASTER_PORT_6379_TCP_PROTO=tcp
REDIS_MASTER_PORT_6379_TCP_PORT=6379
REDIS_MASTER_PORT_6379_TCP_ADDR=10.0.0.11
```

# Service Discovery

## \* DNS

- An optional (though strongly recommended) cluster add-on is a DNS server
- The DNS server watches the Kubernetes API for new Services and creates a set of DNS records for each

For example, if you have a Service called "**redis-master**" in Kubernetes Namespace "**my-ns**"

a DNS record for "**redis-master.my-ns**" is created.

- Kubernetes also supports DNS SRV (service) records for named ports

You can do a **DNS SRV** query for "**\_http.\_tcp.redis-master.my-ns**" to discover the port number for "**http**"

**Load balance**

# Data management

# Persistent Volume

```
kubectl -s http://{HOST}:{PORT} create -f k8s/pv/efs.yml
```

```
cat k8s/pv/efs.yml
```

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv0001
spec:
  capacity:
    storage: 50Gi
  accessModes:
    - ReadWriteOnce
  persistentVolumeReclaimPolicy: Recycle
  nfs:
    path: /
    server: us-west-2a.....efs.us-west-2.amazonaws.com
```

# Persistent Volume Claim

```
kubectl -s http://{HOST}:{PORT} create -f k8s/pvc/efs.yml
```

```
cat k8s/pvc/efs.yml
```

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: efs
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 20Gi
```

# Usage


```
apiVersion: v1
kind: ReplicationController
metadata:
  name: nginx
spec:
  replicas: 1
  selector:
    app: nginx
  template:
    metadata:
      name: nginx
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx
          ports:
            - containerPort: 80
          volumeMounts:
            - mountPath: /data/www
              name: nfs-storage
      volumes:
        - name: nfs-storage
          persistentVolumeClaim:
            claimName: ebs
```



# Health Checking

# Monitoring

`http://{HOST}:{PORT}/ui/`



# Thank you

Luciano Faustino , Tiago Katcipsis

[lucianoborguetti@gmail.com](mailto:lucianoborguetti@gmail.com) , [tiagokatcipsis@gmail.com](mailto:tiagokatcipsis@gmail.com)

(mailto:lucianoborguetti@gmail.com%20%20%20%20%20,%20tiagokatcipsis@gmail.com)

<https://github.com/lborguetti> , <https://github.com/katcipsis>

(https://github.com/lborguetti%20%20,%20https://github.com/katcipsis)

